

## ***The Well-Tempered Semaphore: Theme with Variations***

Kenneth A. Reek, Professor  
Department of Computer Science  
Rochester Institute of Technology  
Rochester, NY 14623

kar@cs.rit.edu

*Rochester Institute of Technology  
Department of Computer Science*


## *Introduction*

Tell me and I know.  
Show me and I remember.  
Let me do it and I understand.


*Confucius*

- Synchronization is difficult!
- Assigning real synchronization problems should improve student comprehension.


*It Looks Easy Enough ...*



## *The Reality—Textbook Interfaces*



## *The Reality—Posix Interface*



```
sem_t S;
```

```
int sem_init( sem_t *S, int pshared,  
              unsigned int value );
```


```
int sem_wait( sem_t *S );  
(But if sem_wait returns -1  
and errno == EINTR,  
you MUST call it again!)
```

```
int sem_post( sem_t *S );
```

```
int sem_destroy( sem_t *S );
```

- SYS V semaphores are worse

## *The Reality—Dijkstra's Semantics*



P( s ):


```
s = s - 1  
if s < 0 then  
    wait on s  
endif
```

V( s ):

```
s = s + 1  
if s <= 0 then  
    unblock one process waiting on s  
endif
```

An unblocked process has priority over newly arriving processes

## *The Reality—Posix Semantics*



```
sem_wait( s ):
    lock mutex
    while count <= 0 do
        unlock mutex
        sleep
        lock mutex
    endwhile
    count = count - 1
    unlock mutex

sem_post( s ):
    lock mutex
    count = count + 1
    unlock mutex
    wake up one sleeping process
```

An unblocked process competes with newly arriving processes

## ***The Semantic Difference***

Process A:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

Process B:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```



## *Posix Semantics*

Process A:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

***A is here***

Process B:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

***B is blocked in the P***

## Posix Semantics

Process A:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

Process B:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

***A leaves the CS***

***B is awakened but not yet running***

- If A continues to run, it can sneak back into the CS before B ever resumes executing

## *Dijkstra's Semantics*

Process A:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

***A is blocked next time around***

Process B:

```
while true do
  noncritical section
  P( s )
  critical section
  V( s )
endwhile
```

***B will continue***

- A cannot sneak back into the CS before B
- In fact, B must be viewed as entering the CS as soon as the V is done.

## ***The Semantic Difference***

- With Posix semantics, one process could indefinitely postpone the other.
- Dijkstra's semantics guarantee that this does not happen.

Therefore:

- Posix semantics are **weak** [Stark]
- Dijkstra's semantics are **strong**

However:

- With three or more processes, any two of them can indefinitely postpone the others
- Neither Dijkstra nor Posix specify the order in which sleeping processes are unblocked (though some implementations do).

## ***Monitors***

A higher-level synchronization technique that ensures mutual exclusion.

- ***wait( condition )*** always blocks the process doing it
- ***signal( condition )*** unblocks one waiting process (ignored if there are none)

## ***Mutual Exclusion in a Monitor***

To maintain mutual exclusion, signal must temporarily delay one process until the other leaves or waits.

- Hoare's semantics: the process doing the signal is delayed, the one that received the signal runs first.
- Brinch Hansen's semantics: the process that received the signal is delayed, the one doing the signal continues.

## *The Semantic Difference*

With Hoare's semantics, execution is constantly switching from the signalling process to the signalled one.


- This code looks bad but may be perfectly good:

```
while count > 0 do  
    signal( condition )  
endwhile
```

Brinch Hansen's semantics are more ``natural``

- The running process keeps running until it blocks itself or leaves

***This Looks Pretty Easy, Too ...***





## ***Monitors and Semaphores***

Monitors are built from semaphores. Will using strong vs. weak semaphores make any difference?

- Let's see ...

## ***A Monitor to Enforce Alternation***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

***Brinch Hansen's Wait and Signal***

```
wait( x ):  
    x.count = x.count + 1  
    if urgentcount > 0 then  
        urgentcount = urgentcount - 1  
        V( urgent )  
    else  
        V( gate )  
    endif  
    P( x.semaphore )  
    P( urgent )  
  
signal( x ):  
    if x.count > 0 then  
        x.count = x.count - 1  
        urgentcount = urgentcount + 1  
        V( x.semaphore )  
    endif
```

***Process A calls alternate***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

*other\_guy.count: 0*

*other\_guy.semaphore: 0*

*urgentcount: 0*

*urgent: 0*

wait( x ):

    x.count = x.count + 1

    if urgentcount > 0 then

        urgentcount = urgentcount - 1

        V( urgent )

    else

        V( gate )

    endif

    P( x.semaphore )

    P( urgent )

signal( x ):

    if x.count > 0 then

        x.count = x.count - 1

        urgentcount = urgentcount + 1

        V( x.semaphore )

    endif

## ***Process B calls alternate: signal***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

*other\_guy.count: ~~0~~ 1*

*other\_guy.semaphore: 0*

*urgentcount: 0*

*urgent: 0*

wait( x ):

    x.count = x.count + 1

    if urgentcount > 0 then

        urgentcount = urgentcount - 1

        V( urgent )

    else

        V( gate )

    endif

    P( x.semaphore ) *← A blocked*

    P( urgent )

signal( x ):

    if x.count > 0 then

        x.count = x.count - 1

        urgentcount = urgentcount + 1

        V( x.semaphore )

    endif

## ***Process B calls alternate: wait (1st part)***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

*other\_guy.count: ~~0~~ 0*

*other\_guy.semaphore: ~~0~~ 1*

*urgentcount: ~~0~~ 1*

*urgent: 0*

wait( x ):

    x.count = x.count + 1

    if urgentcount > 0 then

        urgentcount = urgentcount - 1

        V( urgent )

    else

        V( gate )

    endif

    P( x.semaphore ) ← *A awakened*

    P( urgent )

signal( x ):

    if x.count > 0 then

        x.count = x.count - 1

        urgentcount = urgentcount + 1

        V( x.semaphore )

    endif

## ***Process B calls alternate: wait (2nd part)***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

other\_guy.count: ~~0~~ ~~1~~ ~~0~~ 1

other\_guy.semaphore: ~~0~~ 1

urgentcount: ~~0~~ ~~1~~ 0

urgent: ~~0~~ 1

wait( x ):

    x.count = x.count + 1

    if urgentcount > 0 then

        urgentcount = urgentcount - 1

        V( urgent )

    else

        V( gate )

    endif

    P( x.semaphore ) ← *B is here*

    P( urgent ) ← *A awakened*

signal( x ):

    if x.count > 0 then

        x.count = x.count - 1

        urgentcount = urgentcount + 1

        V( x.semaphore )

    endif

## ***Process B calls alternate: Grand Finale***

condition other\_guy

alternate():

    signal( other\_guy )

    wait( other\_guy )

*other\_guy.count: ~~1~~ ~~1~~ ~~1~~ 1*

*other\_guy.semaphore: ~~1~~ ~~1~~ 0*

*urgentcount: ~~1~~ ~~1~~ 0*

*urgent: ~~1~~ ~~1~~ 0*

wait( x ):

    x.count = x.count + 1

    if urgentcount > 0 then

        urgentcount = urgentcount - 1

        V( urgent )

    else

        V( gate )

    endif

    P( x.semaphore ) *← A reblocked!*

    P( urgent ) *← B departs!*

signal( x ):

    if x.count > 0 then

        x.count = x.count - 1

        urgentcount = urgentcount + 1

        V( x.semaphore )

    endif



## *Analysis*

- A was blocked when B called alternate
- B *should have* blocked, A *should have* continued
- The opposite happened ...
- Conclusion: *IT DIDN'T WORK!* (duh)


## ***Monitors and Semaphores***

Does using strong vs. weak semaphores make any difference?

- Brinch Hansen's semantics: more intuitive, but weak semaphores can cause a failure
- Hoare's semantics: less intuitive, but works correctly with either type of semaphore
- Weak semaphores are not as useful as strong semaphores

***Weak semaphores can cause problems  
when implementing synchronization projects!***

## *Solution: Hide Them with a Wrapper*



Pay no attention to the man behind the curtain, for he is the great and powerful p**O**Zix!



## ***Wrapper: Dijkstra's Counting Semaphores***

Object Oriented implementation in C++ enforces semaphore rules:

- The only operations defined are P and V
- Data is private
- Constructor *requires* an initializer

Option to reverse the unblocking order

- Helps find programs that incorrectly depend on the order

Implemented using Posix (or Solaris) condition variables

- Should port easily to other Posix systems
- Might port to Win32 (perhaps using Events)?

## *Using the Wrapper: Semaphores*

```
#include "Semaphore.h"

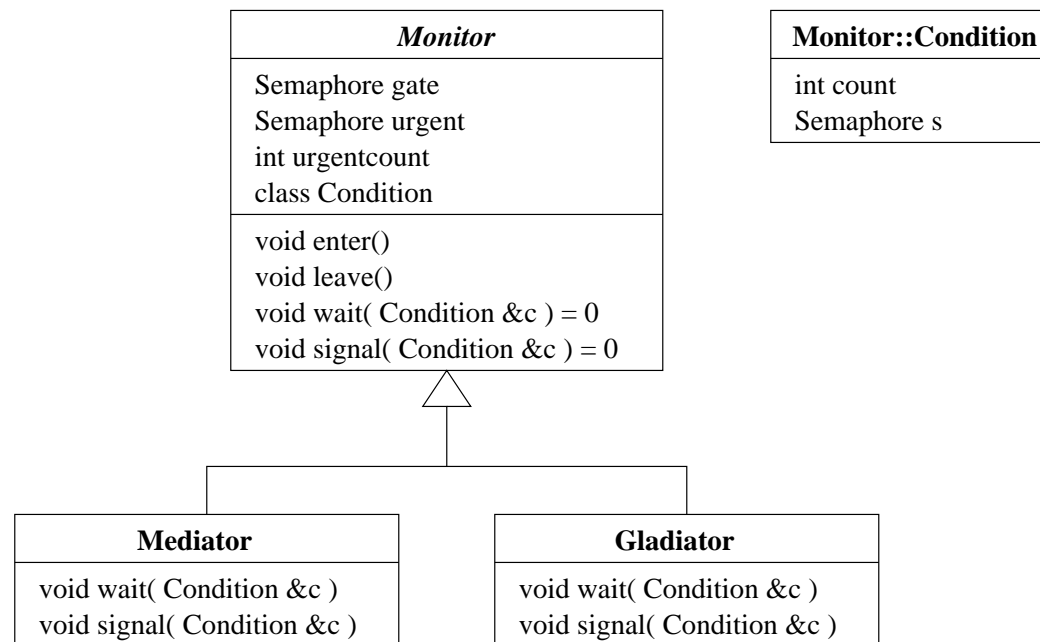
Semaphore mutex( 1 ); // Initialize count to 1

while( true ){
    // non-critical section
    mutex.P();
    // critical section
    mutex.V();
}
```

- Wrapper defines simpler interface to create threads

## Wrapper: Monitors

- Both Hoare's semantics (called Mediator) and Brinch Hansen's semantics (called Gladiator) are implemented



## *Using the Wrapper: Monitors*

```
#include "Monitor.h"

class TwoThreadAlternate: public Gladiator {
public:
    void    alternate();

private:
    Condition other_guy;
};

void TwoThreadAlternate::alternate(){
    enter();           // enter mutual exclusion
    signal( other_guy );
    wait( other_guy );
    leave();           // leave mutual exclusion
}
```

## ***Nifty Assignments: Semaphores***

### Lucky 7's

- Every thread has a number from 1-6
- Incoming threads are blocked until there is a group whose numbers add up to 7; those in the group are unblocked

### Thread groups

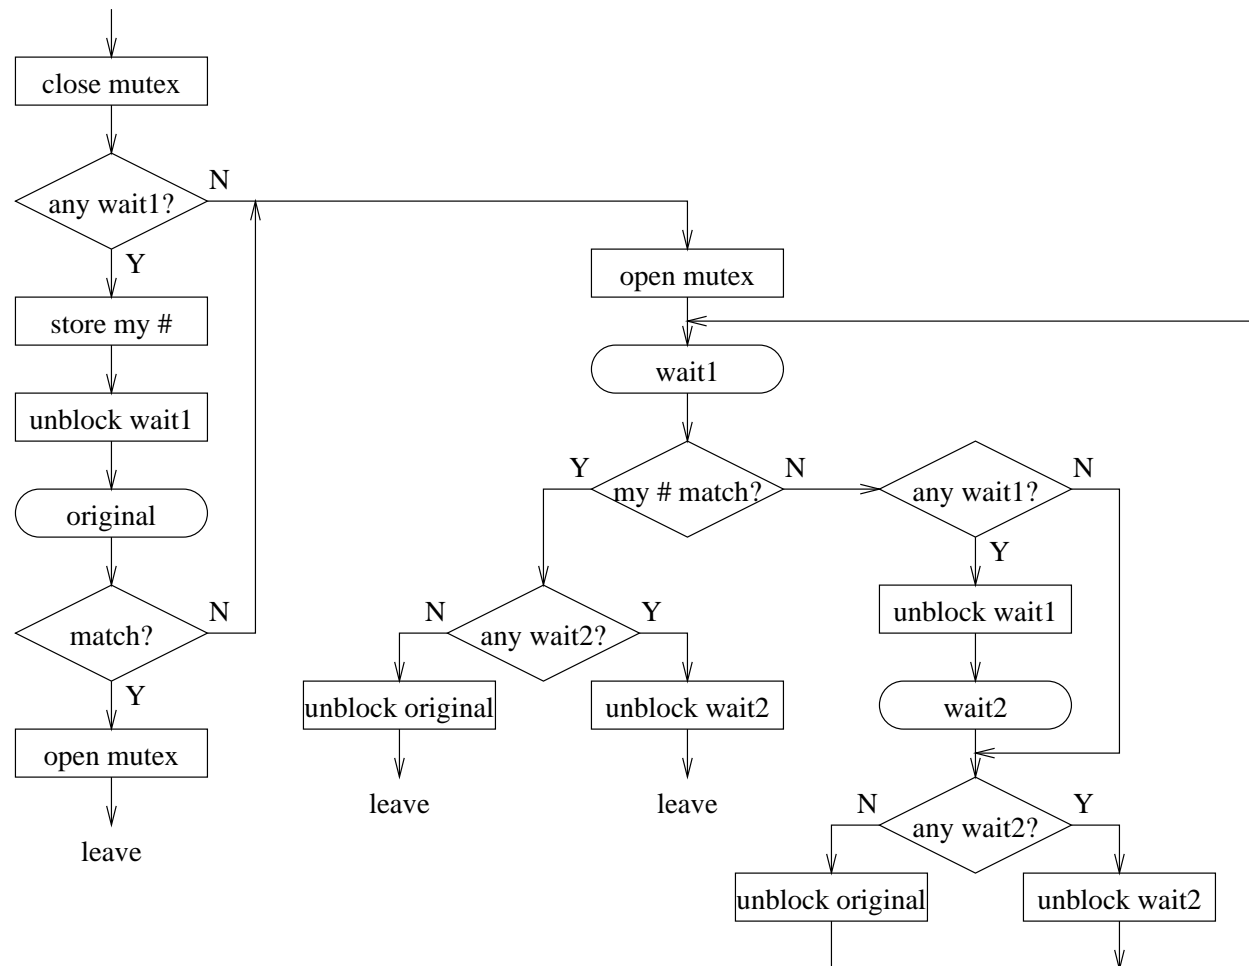
- Incoming threads are blocked until there are enough to form a group, which is then unblocked
- Group size can be changed, which may unblock groups too

### Pairs

- Every thread has an arbitrary integer id
- Each thread is blocked until another with the same id arrives



## *Pairs Logic using Semaphores*



## ***Nifty Assignments: Monitors***

Messages

- Pass a single integer message among threads
- Blocking & non-blocking send, blocking & non-blocking receive

Lucky 7's

Pairs

## ***Pairs Logic using Brinch Hansen's Monitor***

- New arrival saves his ID and unblocks all waiting threads
- Waiting threads compare their ID one by one
- If they don't match, they reblock
- If one matches, it continues
- Last unblocked thread signals the new arrival

***This Stuff is Great! Can We Use It?***

**Sure!**

*But please share any NIFTY ASSIGNMENTS  
you dream up with the rest of us!*

From the World Wide Web:

- Go to [www.cs.rit.edu/~kar](http://www.cs.rit.edu/~kar), click ``*Various papers and colloquia*``

You'll find:

- All of the source code and handouts describing it
- Handouts showing semaphore and monitor solutions to several classic synchronization problems
- The paper and these slides

***Thank you for your attention!***

Kenneth A. Reek, Professor  
Department of Computer Science  
Gollisano College of Computing and Information Sciences  
Rochester Institute of Technology

[www.cs.rit.edu/~kar](http://www.cs.rit.edu/~kar)  
[kar@cs.rit.edu](mailto:kar@cs.rit.edu)